











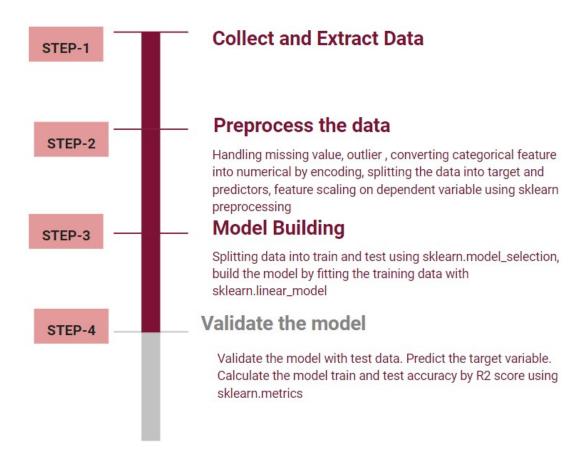




Purpose of the Project

- 1. Exploratory data analysis (EDA) is one of the most significant methods to examine the data we have. In this project, we are going to explore the hidden patterns in the dataset and extract information from them.
- 2. Implementing Multiple Linear Regression on the dataset for future predicion.
- 3. Regularization techniques used to address over-fitting
- 4. Gradient Descent is an optimization algorithm is used that minimize a cost function as far as possible.
- 5. In addition to above, Polynomial regression is used to describes the fitting of a nonlinear relationship if any.
 - In linear regression, the linearity assumption states that there is a linear relationship between the independent variables (also known as predictors or features) and the dependent variable (also known as the target or response variable). This assumption implies that the relationship between the independent variables and the dependent variable can be described by a straight line.
 - Multiple regression suffers from multicollinearity, autocorrelation, heteroskedasticity.
 - Linear Regression is very sensitive to Outliers. It can terribly affect the regression line and eventually the forecasted values.

work flow of the model:



About Dataset

About

Dataset about 50 Startups' expenditures & profits

Column Description

50 startup dataset with columns

- R&D Spend
- Administration
- Marketing Spend
- State
- Profit

Importing Neccessary Libraries

```
In [1]: import os
   import numpy as np
   import pandas as pd
   import matplotlib.pyplot as plt
   import seaborn as sns
   sns.set()
   %matplotlib inline
   import warnings
   warnings.filterwarnings('ignore')
```

Sklearn package used for Machine Learning

```
Model : Linear Regression
```

```
In [2]: from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import StandardScaler, LabelEncoder, PolynomialFeatures
    from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet, SGDRegressor
    from sklearn.metrics import r2_score, mean_absolute_percentage_error, mean_squared_error
    from sklearn import metrics
```

Loading and Reading Data

```
In [3]: dataset = '50_Startups.csv'
    df = pd.read_csv(dataset)
    df.head(10)

Out[3]: R&D Spend Administration Marketing Spend State Profit
```

```
471784.10 New York 192261.83
    165349.20
                    136897.80
    162597.70
                    151377.59
                                     443898.53 California 191792.06
                                                  Florida 191050.39
2
    153441.51
                    101145.55
                                     407934.54
3
    144372.41
                    118671.85
                                     383199.62 New York 182901.99
    142107.34
                    91391.77
                                     366168.42
                                                  Florida 166187.94
    131876.90
                    99814.71
                                     362861.36 New York 156991.12
6
    134615.46
                    147198.87
                                     127716.82 California 156122.51
    130298.13
                    145530.06
                                     323876.68
                                                  Florida 155752.60
    120542.52
                    148718.95
                                     311613.29 New York 152211.77
8
    123334.88
                    108679.17
                                     304981.62 California 149759.96
```

```
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
                  Non-Null Count Dtype
#
   Column
                   -----
    -----
               50 non-null
0
   R&D Spend
                                 float64
                                 float64
    Administration 50 non-null
                                 float64
2
    Marketing Spend 50 non-null
                                 object
    State
                   50 non-null
4 Profit
```

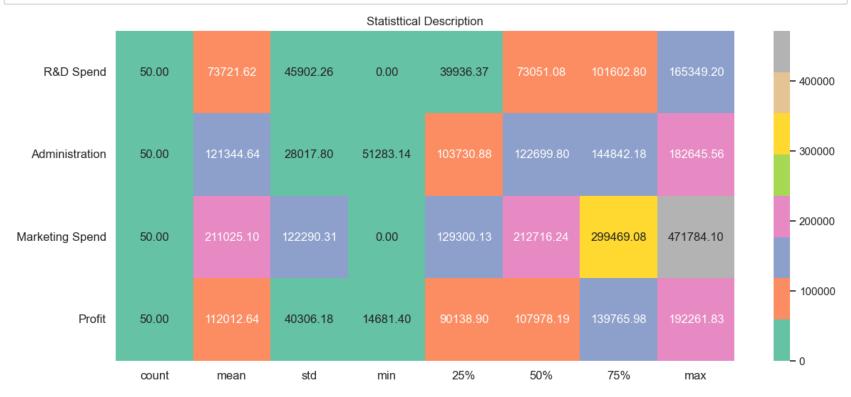
4 Profit 50 non-null dtypes: float64(4), object(1) memory usage: 2.1+ KB

```
In [5]: len(df.columns)
Out[5]: 5
```

```
In [6]: df.shape
Out[6]: (50, 5)
```

Descriptive Statistics (Univariate Analysis)

	CO	uni	mean	Stu	min	25%	50%	15%	illax
R&D Sp	end 5	50.0	73721.6156	45902.256482	0.00	39936.3700	73051.080	101602.8000	165349.20
Administra	tion 5	50.0	121344.6396	28017.802755	51283.14	103730.8750	122699.795	144842.1800	182645.56
Marketing Sp	end 5	50.0	211025.0978	122290.310726	0.00	129300.1325	212716.240	299469.0850	471784.10
Р	rofit 5	50.0	112012.6392	40306.180338	14681.40	90138.9025	107978.190	139765.9775	192261.83



Removing Duplicate Rows

```
In [9]: # Remove duplicates

def drop_dup(df):
    if df.duplicated().any() == True:
        print('The total duplicate row before removing duplicate:', df.duplicated().sum())
        df.drop_duplicates(inplace=True , keep = 'last') # Remove duplicates
        df = df.reset_index(drop=True) #Reset the index
        print('The total duplicate row after removing duplicate:', df.duplicated().sum(), ' \nshape of datase
        else:
            return 'No duplicate entries'
        drop_dup(df)
```

Out[9]: 'No duplicate entries'

Checking Null Value

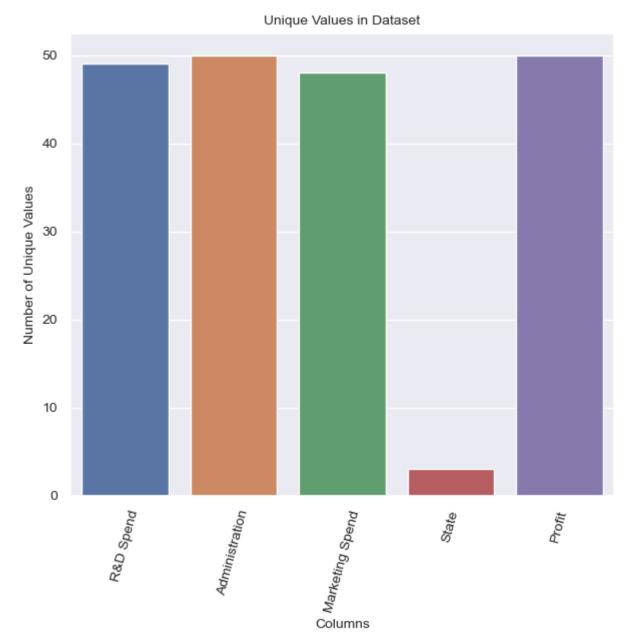
Out[10]:

	col_name	col_type	null_col(%)
0	R&D Spend	float64	0.0
1	Administration	float64	0.0
2	Marketing Spend	float64	0.0
3	State	object	0.0
4	Profit	float64	0.0

dtype: int64

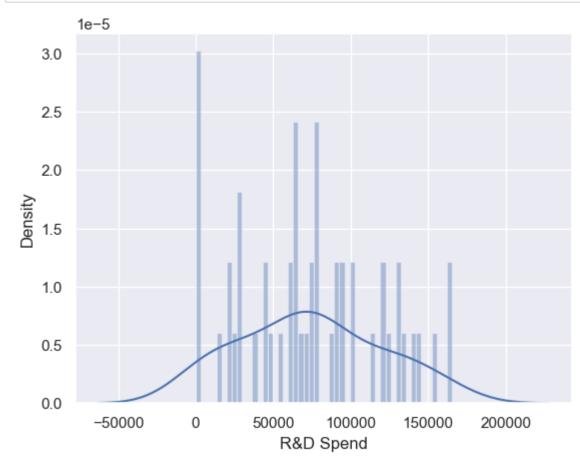
Exploratory Data Analysis (EDA) 📈 📊 🤽

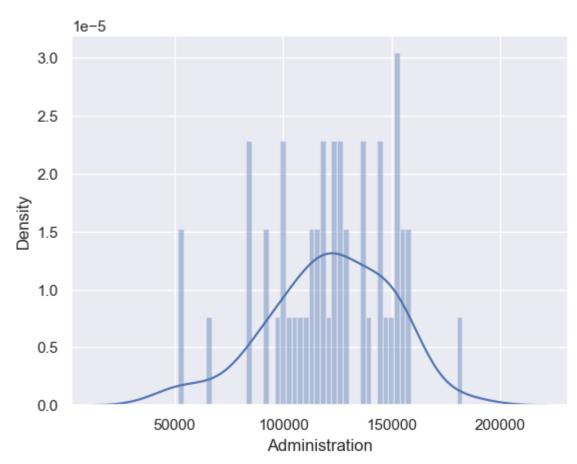
```
In [11]: # Count the number of unique values in each column
         def check unquie count(df):
             unique counts = df.nunique()
             print(unique_counts)
         # Create a bar plot or count plot of unique values
             plt.figure(figsize=(7, 6))
             sns.barplot(x=unique_counts.index, y=unique_counts.values,)
             plt.xticks(rotation=75, fontsize= 10)
             plt.yticks( fontsize= 10 )
             plt.xlabel('Columns', fontsize=10)
             plt.ylabel('Number of Unique Values', fontsize=10)
             plt.title('Unique Values in Dataset', fontsize=10)
         # Display the plot
             plt.show()
         check_unquie_count(df)
         R&D Spend
         Administration
                             50
                             48
         Marketing Spend
         State
                             3
         Profit
                             50
```

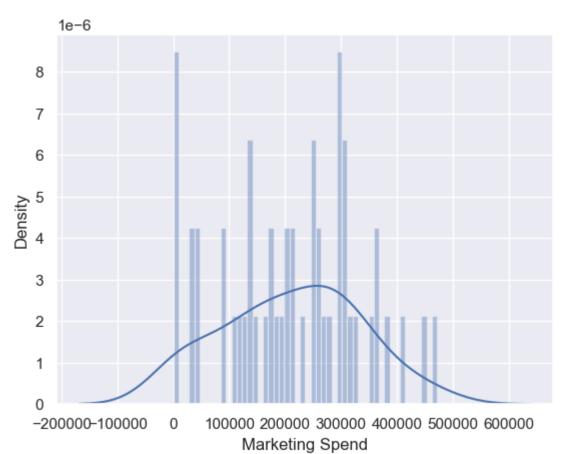


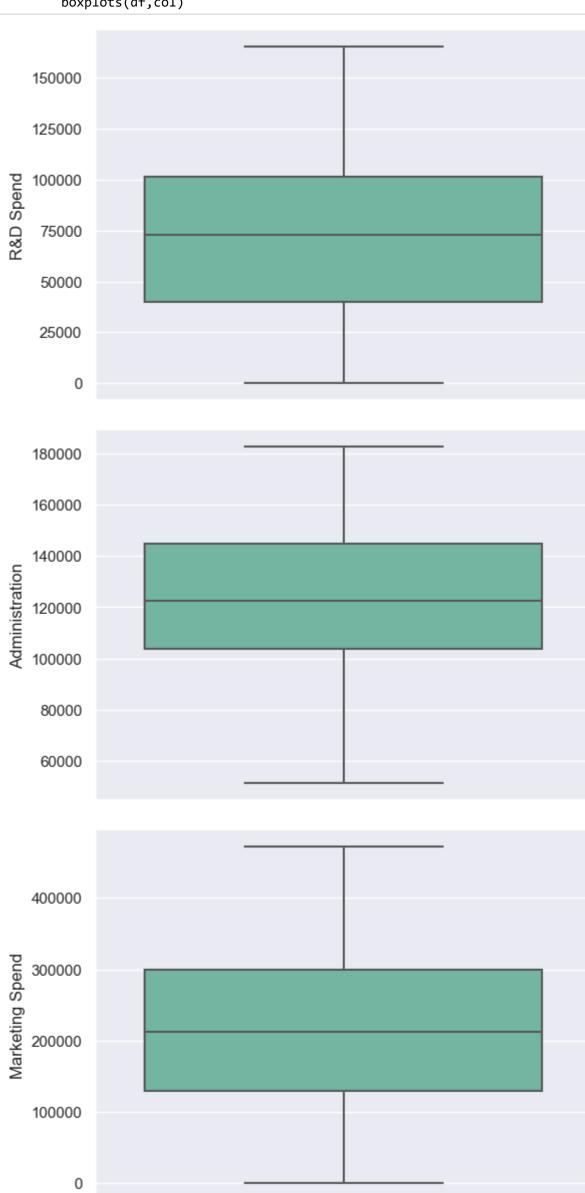
```
In [12]: | num_not_cat = df.select_dtypes(exclude='object')
         num_not_cat.dtypes
Out[12]: R&D Spend
                            float64
                            float64
         Administration
                            float64
         Marketing Spend
         Profit
                            float64
         dtype: object
In [13]: # outlier check on all independent feature except dependent variable
         outlier_list = list(num_not_cat.columns)
         list_remove=['Profit']
         for i in list_remove:
             outlier_list.remove(i)
         outlier_list
Out[13]: ['R&D Spend', 'Administration', 'Marketing Spend']
```

Outlier Check





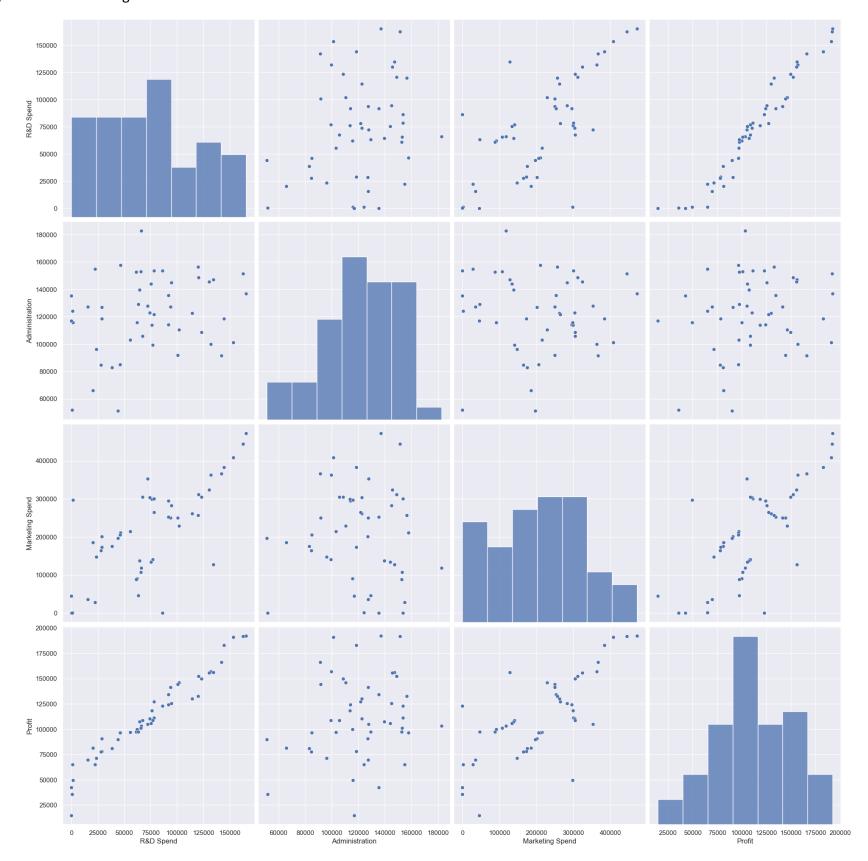




Bivariate Analysis

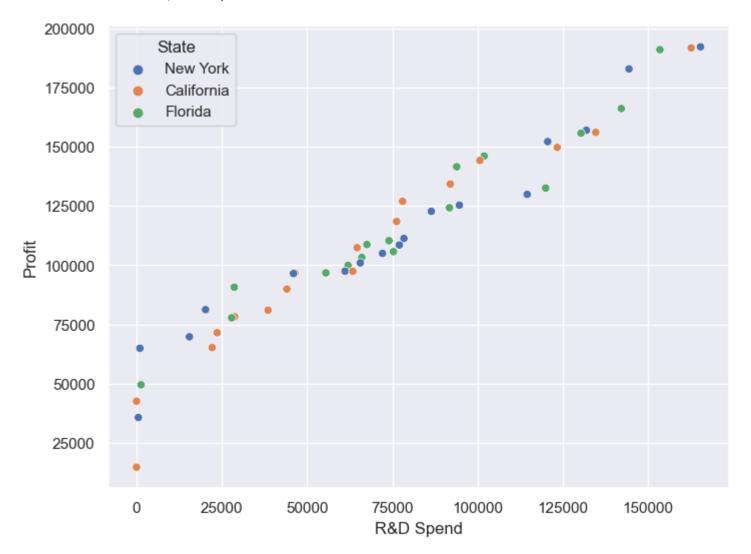
In [16]: sns.pairplot(df, size = 5, kind = 'scatter')

Out[16]: <seaborn.axisgrid.PairGrid at 0x1b2b1fa89d0>



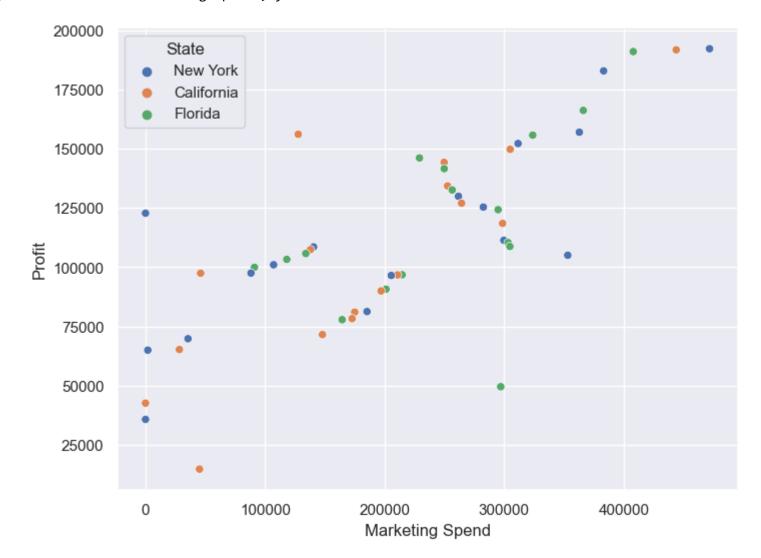
```
In [17]: plt.figure(figsize = (8,6))
sns.scatterplot(x = df['R&D Spend'],y =df['Profit'],hue=df['State'] )
```

Out[17]: <Axes: xlabel='R&D Spend', ylabel='Profit'>



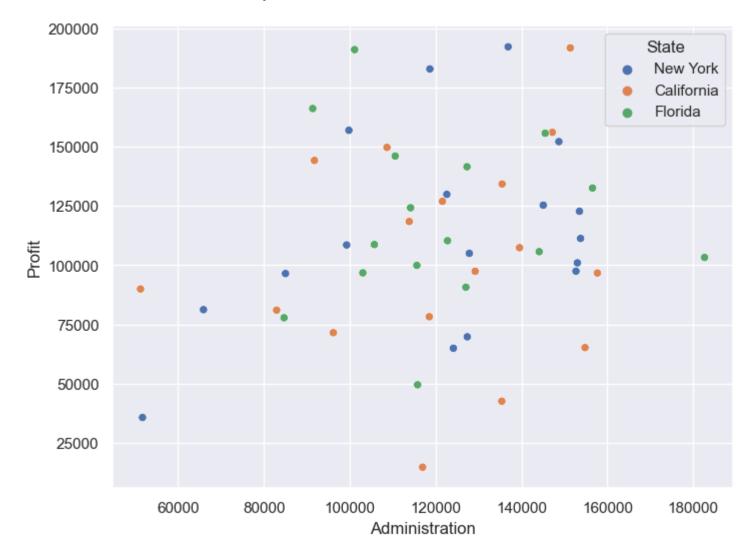
```
In [18]: plt.figure(figsize = (8,6))
sns.scatterplot(x = df['Marketing Spend'],y =df['Profit'],hue=df['State'] )
```

Out[18]: <Axes: xlabel='Marketing Spend', ylabel='Profit'>



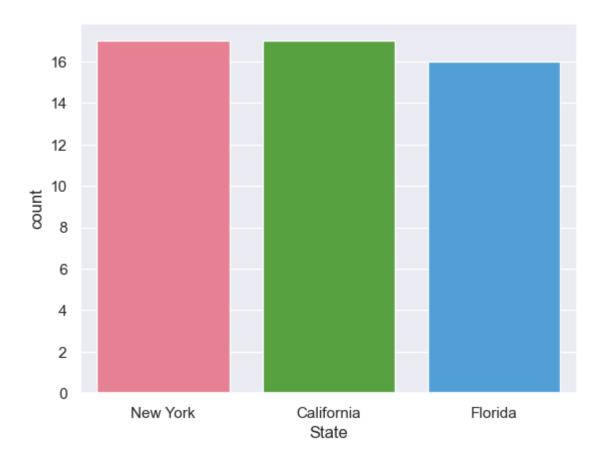
```
In [19]: plt.figure(figsize = (8,6))
sns.scatterplot(x = df['Administration'],y =df['Profit'],hue=df['State'] )
```

Out[19]: <Axes: xlabel='Administration', ylabel='Profit'>



```
In [20]: sns.countplot(x = "State", data = df, palette = 'husl', )
```

Out[20]: <Axes: xlabel='State', ylabel='count'>



Preaparation of Data before Training the algorithm

To train our regression mode, the first step is to split up our data into an target(dependent variables) and features (predictors). And we need to do feature selection only on the predictors or independent variable.

Step-1: Spliting into dependent and independent variable

```
In [21]: list(df)
Out[21]: ['R&D Spend', 'Administration', 'Marketing Spend', 'State', 'Profit']
```

```
In [22]: # Assuming 'y' is the column name of the target variable
          target = 'Profit'
          y = df[[target]]
          # Assuming 'X' is the DataFrame containing the feature columns
          features = df.drop(target, axis=1)
          x = features
In [23]: x.head()
Out[23]:
              R&D Spend Administration Marketing Spend
                                                         State
           0
              165349.20
                             136897.80
                                            471784.10 New York
               162597.70
                             151377.59
                                            443898.53 California
                             101145.55
           2
               153441.51
                                            407934.54
                                                        Florida
           3
               144372.41
                             118671.85
                                            383199.62 New York
               142107.34
                              91391.77
                                            366168.42
                                                        Florida
```

```
In [24]: x.shape
```

Out[24]: (50, 4)

In [25]: y.head()

Out[25]:

Profit

- **0** 192261.83
- **1** 191792.06
- **2** 191050.39
- **3** 182901.99
- **4** 166187.94

Step-2: Encoding of categorical feature

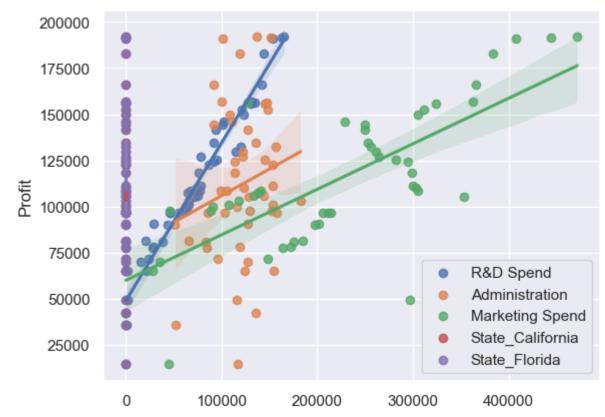
'State' feature is object type. In principle, the object type data are treated under encoding after converting the, into category.

```
In [26]: for col in x.columns:
    if x[col].dtypes == 'object' and x[col].nunique() > 2:
        x[col].astype('category')
        x = pd.get_dummies(x, columns = [col], drop_first = True)

#x = x.drop(['State_New York' ], axis = 1)
    x.head()
```

Out[26]:

	R&D Spend	Administration	Marketing Spend	State_California	State_Florida
0	165349.20	136897.80	471784.10	0	0
1	162597.70	151377.59	443898.53	1	0
2	153441.51	101145.55	407934.54	0	1
3	144372.41	118671.85	383199.62	0	0
4	142107.34	91391.77	366168.42	0	1

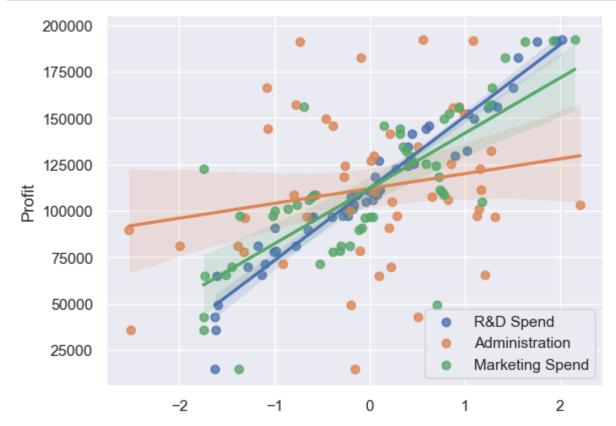


step -3: Feature Scaling

• Feature scaling is done only on dependent variables. So we need to split the data into target and dependent variable.

```
In [28]: sc = StandardScaler()
          sc_x = sc.fit_transform(x)
          sc_x = pd.DataFrame(sc_x)
In [29]: sc_x.head()
Out[29]:
                            1
                                     2
                                               3
                                                        4
           0 2.016411 0.560753 2.153943 -0.717741 -0.685994
           1 1.955860
                      1.082807 1.923600 1.393261 -0.685994
           2 1.754364 -0.728257 1.626528 -0.717741
                                                  1.457738
           3 1.554784 -0.096365 1.422210 -0.717741 -0.685994
           4 1.504937 -1.079919 1.281528 -0.717741 1.457738
In [30]: list(x)
Out[30]: ['R&D Spend',
           'Administration',
           'Marketing Spend',
           'State_California',
           'State_Florida']
```

```
In [31]: sns.regplot(x=sc_x[0], y=y, label='R&D Spend',)
    sns.regplot(x=sc_x[1], y=y, label='Administration',)
    sns.regplot(x=sc_x[2], y=y, label='Marketing Spend',)
    plt.xlabel('')
    plt.legend()
    plt.show()
```



step -4: Multi-collinearity check

There should be no or little mulyicollinearity present for the model building



VIF (Multi-collinearity Check)

Multicollinearity refers to a situation in regression analysis where there is a high correlation between two or more predictor variables (also known as independent variables or features). It occurs when the predictor variables in a regression model are highly linearly related to each other, making it difficult to distinguish the individual effects of each variable on the dependent variable.

In the presence of multicollinearity, it becomes challenging to determine the true relationship between the predictor variables and the target variable. This is because multicollinearity can lead to unstable and unreliable estimates of the regression coefficients.

Identifying multicollinearity: Common methods to detect multicollinearity include calculating correlation matrices, variance inflation factors (VIF> 5).

Dealing with multicollinearity: If multicollinearity is detected, several strategies can be employed,

• Dropping one or more correlated variables from the model. (incase more than 1 features have VIF > 5, drop highest and need to perform VIF check again.

```
In [34]: vif
```

Out[34]:

Features	VIF	
R&D Spend	2.495511	0
Administration	1.177766	1
Marketing Spend	2.416797	2
State_California	1.335061	3
State Florida	1.361299	4

• for VIF > 5 refers to multicollinearity. For linear regression model building we need no or little multicollinearitywe need to drop the feature. As we don't need multi-collinearity

step-5: Split the data for building the model and prediction

**CASE- I : Test size = 0.3 **

Split the data (70% Train and 30% Test)

```
In [35]: # Split the data into training and test for building the model and for prediction
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=121)
    print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)
    (35, 5) (15, 5) (35, 1) (15, 1)
```

Approach no - 1: Multiple LINEAR REGRESSION 📈 📉

Multiple linear equation;

$$\hat{y} = y + \varepsilon$$

$$\hat{y} = \beta_0 + \beta_1 \cdot x_1 + \beta_2 \cdot x_2 + \dots + \beta_n \cdot x_n + \varepsilon$$

$$\hat{y} = \beta_0 + \Sigma_1^i \beta_i \cdot x_i + \varepsilon$$

where β_0 is intercept and β_i 's are slopes and ε is error.

- Response = (Constant + Slope * Predictors) + Error
- y is the Response, x_i 's are the Predictors
- The difference between the actual value and the model's estimate a residual (error).
- These residuals will play a significant role in judging the usefulness of a model.
- If the residuals are small, it implies that the model is a good estimator.

```
<img src="hyperplane.jpg" alt="" width="400" height="300">
```

Performance Matrix

Here are three common evaluation metrics for regression problems:

• Mean Absolute Error (MAE) is the mean of the absolute value of the errors:

$$\frac{1}{n}\sum_{i=1}^n|y_i-\hat{y}_i|$$

• Mean Squared Error (MSE) is the mean of the squared errors:

$$\frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2$$

• Root Mean Squared Error (RMSE) is the square root of the mean of the squared errors:

$$\sqrt{\frac{1}{n}\sum_{i=1}^{n}(y_i-\hat{y}_i)^2}$$

- ★ Comparing these metrics:
- MAE is the easiest to understand, because it's the average error.
- MSE is more popular than MAE, because MSE "punishes" larger errors, which tends to be useful in the real world.
- RMSE is even more popular than MSE, because RMSE is interpretable in the "y" units.

All of these are **loss functions**, because we want to minimize them.

R^2 Score

$$R^{2} = 1 - \frac{SSR}{TSS} = 1 - \frac{\sum_{i} (y_{i} - \hat{y}_{i})^{2}}{\sum_{i} (y_{i} - \bar{y})^{2}}$$

- TSS: Total Sum of Square = $\sum_{i} (Actual)_{i} (MeanActual) = \sum_{i} (y_{i} \bar{y})^{2}$
- SSR : Sum of Square residual = $(Residual)^2 = (Actual Predicted)^2 = \sum_i (y_i \hat{y}_i)^2$

Adjusted R^2

$$R_{adj}^2 = 1 - \frac{(1-R^2)(N-1)}{(N-p-1)}$$

- N = Total sample size
- p = Number of predictors (No. of Independent variable)

Training

```
In [36]: # Tain the model with LR model
LR = LinearRegression()
LR.fit(x_train, y_train)
```

Out[36]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [37]: # find the slopes and intercept from the trained model
   intercept =LR.intercept_
   slope =LR.coef_

   print('The intercept for our linear model is :',intercept.round(3))
   print()
   print('The coefficients are :',slope.round(3))
```

The intercept for our linear model is : [43887.251]

The coefficients are : [[8.480000e-01 -1.900000e-02 1.800000e-02 7.453231e+03 3.063963e+03]]

```
In [39]: Lr_data = {'Slope': slope, 'Features' : x.columns}
         coeff_df = pd.DataFrame(slope, columns=['Slope'], index =x.columns)
         coeff_df
                             Slope
```

Out[39]:

```
R&D Spend
                    0.848180
 Administration
                    -0.019409
Marketing Spend
                    0.018301
State_California 7453.231161
   State_Florida 3063.963122
```

Prediction of target variable using Linear Regression model

```
In [40]: # Predict house price by using linear Regression model with test dataset
         y_pred_train = LR.predict(x_train)
         err_train = y_train - y_pred_train
         Train_accuracy = r2_score(y_train, y_pred_train)
         y_pred_test = LR.predict(x_test)
         err_test = y_test - y_pred_test
         Test_accuracy = r2_score(y_test, y_pred_test)
         print ('Train accuracy :', Train_accuracy,'\n' 'Test accracy :', Test_accuracy)
         Train accuracy : 0.9659473642911475
         Test accracy : 0.8990798108840938
In [41]: err train.skew()
Out[41]: Profit -0.306667
         dtype: float64
In [42]: err_train.kurtosis()
Out[42]: Profit
                   0.42166
         dtype: float64
```

According to the histogram the error terms are normally distributed and that is then further confirmed by skewness and kurtosis which are close to zero.

```
In [43]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
         # Plot regression for training data
         for i in range(len(x_train.columns)):
             sns.regplot(x=x_train.iloc[:,i ], y=y_pred_train, ax=ax1, )
             ax1.set_title('Training data')
             ax1.set_xlabel('X values')
         # Plot regression for test data
         for i in range(len(x_test.columns)):
             sns.regplot(x=x_test.iloc[:, i], y=y_pred_test, ax=ax2)
             ax2.set_title('Test data')
             ax2.set xlabel('X values')
         # Adjust spacing between subplots
         plt.tight_layout()
         # Show the plot
         plt.show()
```



```
In [44]: # mse, rmse, mape
         MSE = mean_squared_error(y_test,y_pred_test)
         print(f'1. mean squared error (MSE) = ',MSE)
         RMSE = np.sqrt(MSE)
         print(f'2. root mean squared error (RMSE) = {RMSE}')
         MAPE = mean_absolute_percentage_error(y_test,y_pred_test)
         print(f'3. mean absolute percentage error (MAPE) = {MAPE}')
         ACC = 100-MAPE*100
         print(f'4. accuracy of the model = {ACC}')
         1. mean squared error (MSE) = 194052164.52818343
```

- 2. root mean squared error (RMSE) = 13930.260748750665
- 3. mean absolute percentage error (MAPE) = 0.25548939871234455
- 4. accuracy of the model = 74.45106012876555

Approach no 2 : Ordinary Least Square (OLS) Method

```
In [46]: | from statsmodels.regression.linear_model import OLS
           import statsmodels.regression.linear_model as smf
           reg_model = smf.OLS(endog = y_train, exog=x_train).fit() # with unscaled data
In [47]:
           reg_model.summary()
Out[47]:
           OLS Regression Results
                Dep. Variable:
                                        Profit
                                                                             0.993
                                                  R-squared (uncentered):
                                                                             0.992
                      Model:
                                         OLS Adj. R-squared (uncentered):
                     Method:
                                Least Squares
                                                               F-statistic:
                                                                             854.0
                       Date: Mon, 17 Jul 2023
                                                         Prob (F-statistic): 2.20e-31
                       Time:
                                     22:16:10
                                                          Log-Likelihood:
                                                                           -373.19
            No. Observations:
                                                                    AIC:
                                          35
                                                                             756.4
                                          30
                                                                    BIC:
                                                                             764.2
                Df Residuals:
                    Df Model:
                                           5
             Covariance Type:
                                    nonrobust
                                  coef
                                          std err
                                                      t P>|t|
                                                                  [0.025
                                                                            0.975]
                 R&D Spend
                                0.8442
                                           0.065 12.900 0.000
                                                                   0.711
                                                                             0.978
              Administration
                                0.2484
                                           0.037
                                                  6.718 0.000
                                                                   0.173
                                                                             0.324
            Marketing Spend
                                0.0429
                                           0.022
                                                  1.932 0.063
                                                                  -0.002
                                                                             0.088
             State_California
                            1.253e+04 4782.658
                                                  2.620 0.014 2761.040 2.23e+04
                                                                -901.563 1.73e+04
                            8205.2066 4459.135
               State_Florida
                                                  1.840 0.076
                             0.833
                                     Durbin-Watson:
                 Omnibus:
                                                        2.004
            Prob(Omnibus):
                            0.659 Jarque-Bera (JB):
                                                        0.362
                                          Prob(JB):
                     Skew: -0.245
                                                        0.834
                  Kurtosis: 3.093
                                          Cond. No. 8.71e+05
```

Notes:

- [1] R² is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The condition number is large, 8.71e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [48]: reg_model = smf.OLS(endog = y_test, exog=x_test).fit() # with unscaled data
reg_model.summary()

Out[48]:

OLS Regression Results

Dep. Variable:	:	Profit	R-squared (uncentered):		0.983	
Model:	:	OLS A	Adj. R-squared (uncentered):		0.975	
Method:	Least So	quares	F-statistic:		F-statistic:	115.8
Date:	: Mon, 17 Ju	ıl 2023	Prob (F-statistic):		-statistic):	1.62e-08
Time	: 22	::16:10	Log-Likelihood:		ikelihood:	-164.38
No. Observations:	:	15			AIC:	338.8
Df Residuals:	:	10			BIC:	342.3
Df Model:	:	5				
Covariance Type:	non	robust				
	coef	std err	t	P> t	[0.025	0.975]
R&D Spend	0.5366	0.155		0.006	0.191	0.882
Administration	0.4017	0.081		0.001	0.221	0.583
Marketing Spend	0.1404	0.054	2.602	0.026	0.020	0.261
State California	-5149.8625	9453.054		0.598	-2.62e+04	1.59e+04
– State Florida	-4237.5460	1.47e+04		0.779	-3.7e+04	2.86e+04
_						
Omnibus:	4.375 D u	ırbin-Wats	on:	2.350		
Prob(Omnibus):	0.112 Jar q	jue-Bera (J	IB):	2.026		
Skew:	-0.840	Prob(J	IB):	0.363		
Kurtosis:	3.645	Cond.	No. 8.55	5e+05		


Notes

- [1] R² is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The condition number is large, 8.55e+05. This might indicate that there are strong multicollinearity or other numerical problems.
 - Endoginity Problem : if dependent variable is set wrongly that is endoginity problem

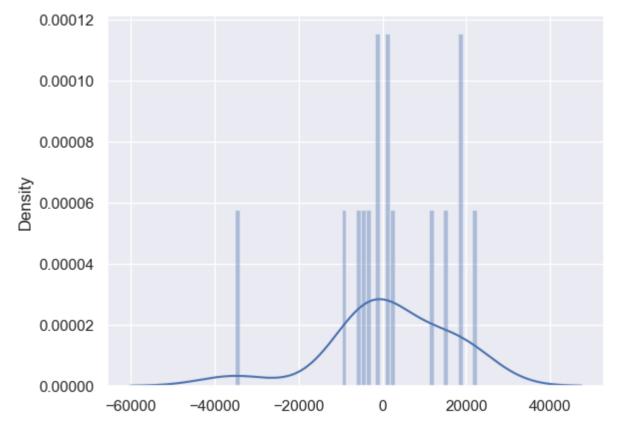
In [49]: # Check linearity

plt.scatter(y_test, y_pred_test)

Out[49]: <matplotlib.collections.PathCollection at 0x1b2b7bb4d30>



```
In [50]: # Normality of Residual
sns.distplot((y_test - y_pred_test), bins=50)
plt.show()
```



Assumption of LR Model is satisfied

- 1) Linearity Satisfied
- 2) Normality of Residuals-Satisfied
- 3) Homoscedasticity Satisfied (there is no outlier and residual is normaly distributed)
- 4) No autocorrelation Satisfied
- 5) No or little Multicollinearity satisfied
- 6) No endogenity problem satisfied

In order to create less complex model when you have a large number of features in your dataset, some of the Regularization techniques used to address over-fitting and feature selection are:

Lasso (L1 reguralization):

Mathematically, it consists of a linear model trained with ℓ_1 prior as regularizer. The objective function to minimize is:

$$\arg\min_{\beta_o,\beta_j} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j|$$

The lasso estimate thus solves the minimization of the least-squares penalty. λ controls the strength of reguralization

Ridge:

Ridge regression addresses some of the problems of **Ordinary Least Squares** by imposing a penalty on the size of coefficients. The ridge coefficients minimize a penalized residual sum of squares,

$$\arg\min_{\beta_o,\beta_j} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2$$

 $\lambda >= 0$ is a complexity parameter that controls the amount of shrinkage: the larger the value of λ , the greater the amount of shrinkage and thus the coefficients become more robust to collinearity.

Ridge regression is an L2 penalized model. Add the squared sum of the weights to the least-squares cost function.

ElasticNet

Test Accuracy: 0.9011734894553887

The first term represents the residual sum of squares, the second term is the L1 regularization term (Lasso), and the third term is the L2 regularization term (Ridge). The Elastic Net regularization combines both L1 and L2 regularization, allowing for a balance between feature selection and coefficient shrinkage. The parameters λ_1 and λ_2 control the strengths of the respective regularization terms.

$$\arg\min_{\beta_o,\beta_j} \sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda_1 \sum_{j=1}^p |\beta_j| + \lambda_2 \sum_{j=1}^p \beta_j^2$$

```
In [53]: elastic = ElasticNet(alpha=0.3, l1_ratio=0.1)
    elastic.fit(x_train, y_train)
    y_pred_train_elastic = elastic.predict(x_train)
    y_pred_test_elastic = elastic.predict(x_test)
    print("Training Accuracy :", r2_score(y_train, y_pred_train_elastic))
    print()
    print("Test Accuracy :", r2_score(y_test, y_pred_test_elastic))
```

Training Accuracy : 0.9632279233912138

Test Accuracy: 0.9179513210700663

Gradient Descent

Gradient Descent is a very generic optimization algorithm capable of finding optimal solutions to a wide range of problems. The general idea of Gradient Sescent is to tweak parameters iteratively in order to minimize a cost function. Gradient Descent measures the local gradient of the error function with regards to the parameters vector, and it goes in the direction of descending gradient. Once the gradient is zero, you have reached a minimum.

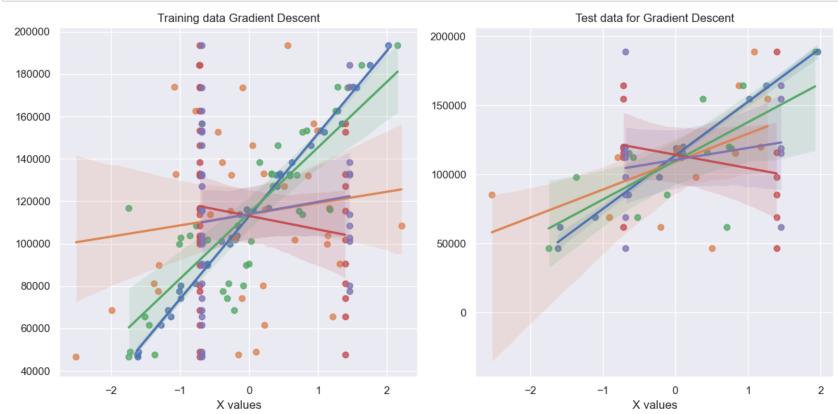
```
In [74]: | from sklearn.model_selection import train_test_split
         x_train, x_test, y_train, y_test = train_test_split(sc_x, y, test_size=0.25, random_state=101)
         print('x_train shape :',x_train.shape,'\n''x_test shape :', x_test.shape,'\n' 'y_train shape :',y_train.shape
         print()
         gd = SGDRegressor()
         gd.fit(x_train, y_train)
         y_pred_gd_train = gd.predict(x_train)
         y_pred_gd_test = gd.predict(x_test)
         Train_accuracy_gd =r2_score(y_train, y_pred_gd_train)
         Test_accuracy_gd = r2_score(y_test, y_pred_gd_test)
         print("GD Trainging Accuracy :", Train_accuracy_gd )
                                      :",Test_accuracy_gd )
         print("GD Test Accuracy
         x_{train} shape : (37, 5)
         x_test shape : (13, 5)
         y_train shape : (37, 1)
         y_test shape : (13, 1)
         GD Trainging Accuracy : 0.9472665189054107
         GD Test Accuracy
                               : 0.9409293485798228
In [55]: | fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
```

```
In [55]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))

# Plot regression for training data
for i in range(len(x_train.columns)):
    sns.regplot(x=x_train.iloc[:,i ], y=y_pred_gd_train, ax=ax1, )
    ax1.set_title('Training data Gradient Descent')
    ax1.set_xlabel('X values')

# Plot regression for test data
for i in range(len(x_test.columns)):
    sns.regplot(x=x_test.iloc[:, i], y=y_pred_gd_test, ax=ax2)
    ax2.set_title('Test data for Gradient Descent')
    ax2.set_xlabel('X values')

# Adjust spacing between subplots
plt.tight_layout()
# Show the plot
plt.show()
```

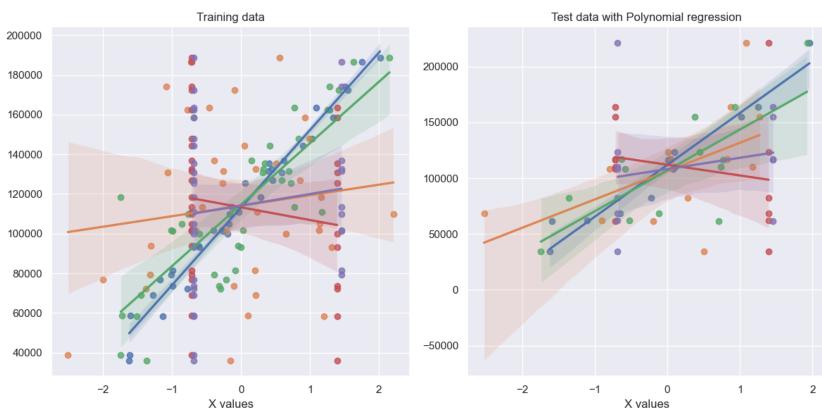


Polynomial regression

```
In [56]: poly = PolynomialFeatures()
    x_train_trans = poly.fit_transform(x_train)
    x_test_trans = poly.fit_transform(x_test)
    lr = LinearRegression()
    lr.fit(x_train_trans, y_train)
    poly.fit(x_train, y_train)
    y_pred_poly_train = lr.predict(x_train_trans)
    y_pred_poly_test = lr.predict(x_test_trans)
    Train_accuracy_poly = r2_score(y_train, y_pred_poly_train)
    Test_accuracy_poly = r2_score(y_test, y_pred_poly_test)

print("Polynomial Regression Trainging Accuracy :", Train_accuracy_poly )
    print("Polynomial Regression Test Accuracy :", Test_accuracy_poly )
```

Polynomial Regression Trainging Accuracy : 0.9643335626660127
Polynomial Regression Test Accuracy : 0.8670425426922718



**CASE-I: Test size = 0.2 **

Split the data (80% Train and 20% Test)

```
In [58]: #Split the data into training and test for building the model and for prediction
         x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=121)
         print('x_train shape :',x_train.shape,'\n''x_test shape :', x_test.shape,'\n' 'y_train shape :',y_train.shape
         y_pred_train = LR.predict(x_train)
         y_pred_test = LR.predict(x_test)
         # Validate the actual price of the test data and predicted price
         Test_accuracy = r2_score(y_test, y_pred_test)
         Train_accuracy = r2_score(y_train, y_pred_train)
         print ('Train accuracy :', Train_accuracy,'\n' 'Test accracy :', Test_accuracy)
         # MSE. RMSE. MAPE
         MSE = mean_squared_error(y_test,y_pred_test)
         print(f'1. mean squared error (MSE) = ',MSE)
         RMSE = np.sqrt(MSE)
         print(f'2. root mean squared error (RMSE) = {RMSE}')
         MAPE = mean_absolute_percentage_error(y_test,y_pred_test)
         print(f'3. mean absolute percentage error (MAPE) = {MAPE}')
         ACC = 100-MAPE*100
         print(f'4. accuracy of the model = {ACC}')
         x_{train} shape : (40, 5)
         x_test shape : (10, 5)
         y_train shape : (40, 1)
         y_test shape : (10, 1)
         Train accuracy: 0.9328286436777007
         Test accracy : 0.9830641471002174
         1. mean squared error (MSE) = 29104086.61012637
         2. root mean squared error (RMSE) = 5394.820350125328
         3. mean absolute percentage error (MAPE) = 0.04542249274560596
         4. accuracy of the model = 95.45775072543941
```

In [59]:

with unscaled data

```
reg_model.summary()
Out[59]:
            OLS Regression Results
                 Dep. Variable:
                                           Profit
                                                      R-squared (uncentered):
                                                                                  0.987
                                            OLS Adj. R-squared (uncentered):
                        Model:
                                                                                  0.985
                                                                   F-statistic:
                      Method:
                                  Least Squares
                                                                                  524.2
                         Date: Mon, 17 Jul 2023
                                                             Prob (F-statistic): 7.50e-32
                                                              Log-Likelihood:
                                        22:16:15
                                                                                -437.43
                         Time:
             No. Observations:
                                             40
                                                                         AIC:
                                                                                  884.9
                 Df Residuals:
                                             35
                                                                         BIC:
                                                                                  893.3
                     Df Model:
                                              5
                                      nonrobust
              Covariance Type:
                                                                      [0.025
                                                                                 0.975]
                                    coef
                                             std err
                                                         t P>|t|
                  R&D Spend
                                              0.078 9.792 0.000
                                                                                 0.918
                                  0.7605
                                                                       0.603
                                  0.3080
                                                                       0.225
               Administration
                                              0.041 7.539 0.000
                                                                                 0.391
```

reg_model = smf.OLS(endog = y_train, exog=x_train).fit()

 Omnibus:
 0.002
 Durbin-Watson:
 1.949

 Prob(Omnibus):
 0.999
 Jarque-Bera (JB):
 0.108

 Skew:
 -0.007
 Prob(JB):
 0.948

 Kurtosis:
 2.746
 Cond. No.
 8.33e+05

0.0633

Notes:

Marketing Spend

State_California

State_Florida

- [1] R² is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

0.028 2.287 0.028

3888.8627 5545.837 0.701 0.488 -7369.785 1.51e+04

4037.2964 5818.962 0.694 0.492 -7775.823

[3] The condition number is large, 8.33e+05. This might indicate that there are strong multicollinearity or other numerical problems.

0.007

0.119

1.59e+04

```
sns.distplot((y_test - y_pred_test), bins=50)
plt.show()
```

- 1) Linearity Satisfied
- 2) Normality of Residuals- Satisfied
- 3) Homoscedasticity Satisfied (there is no outlier and residual is normaly distributed, no variance in residual)
- 4) No autocorrelation Satisfied
- 5) No or little Multicollinearity satisfied
- 6) No endogenity problem satisfied

P-values and coefficients in regression analysis work together to tell which relationships in the model are statistically significant and the nature of those relationships.

- The coefficients describe the mathematical relationship between each independent variable and the dependent variable.
- The p-values for the coefficients indicate whether these relationships are statistically significant.

reg_model.summary()

In [60]:

Out[60]:

with unscaled data

```
OLS Regression Results
     Dep. Variable:
                              Profit
                                         R-squared (uncentered):
                                                                     0.995
                               OLS Adj. R-squared (uncentered):
           Model:
                                                                     0.991
          Method:
                                                                     210.7
                      Least Squares
                                                      F-statistic:
             Date: Mon, 17 Jul 2023
                                               Prob (F-statistic): 8.29e-06
                           22:16:15
                                                 Log-Likelihood:
                                                                   -104.46
            Time:
No. Observations:
                                10
                                                            AIC:
                                                                     218.9
     Df Residuals:
                                                            BIC:
                                                                     220.4
                                 5
                                 5
        Df Model:
 Covariance Type:
                          nonrobust
                                std err
                                                         [0.025
                                                                   0.975]
                        coef
                                            t P>|t|
      R&D Spend
                      0.6219
                                 0.137 4.543 0.006
                                                                    0.974
                                                          0.270
  Administration
                      0.2576
                                 0.112 2.293 0.070
                                                         -0.031
                                                                    0.547
Marketing Spend
                      0.1323
                                 0.041 3.203 0.024
                                                          0.026
                                                                    0.238
 State_California
                   1.27e+04 9393.275 1.352 0.234 -1.14e+04 3.68e+04
                  7324.7987 1.41e+04 0.519 0.626
                                                       -2.9e+04 4.36e+04
    State_Florida
      Omnibus: 1.822
                          Durbin-Watson:
                                              2.054
Prob(Omnibus): 0.402 Jarque-Bera (JB):
                                              0.717
          Skew: 0.652
                                Prob(JB):
                                              0.699
```

reg_model = smf.OLS(endog = y_test, exog=x_test).fit()

Notes:

- [1] R² is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Cond. No. 1.14e+06

[3] The condition number is large, 1.14e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Reguralization

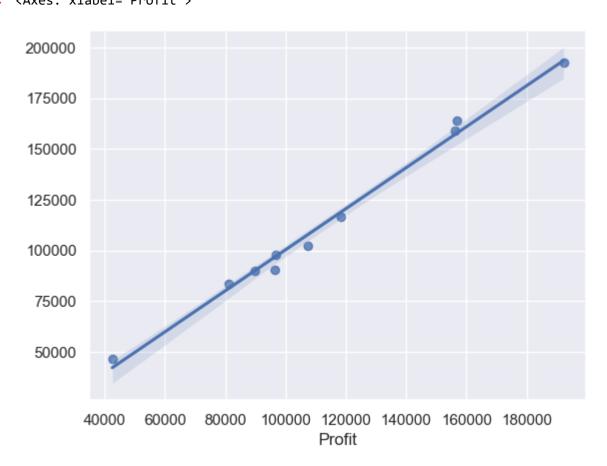
Kurtosis: 2.855

LASSO

Ridge

```
In [62]: # Part 2 : Ridge Regression (L2- Regularization)
         # closure to zero but not exact zero
         # penalty - 0.3
         ridge = Ridge(alpha=0.3)
         ridge.fit(x_train, y_train)
         print("Ridge Model :", (ridge.coef_))
         y_pred_train_ridge = ridge.predict(x_train)
         y_pred_test_ridge = ridge.predict(x_test)
         print("Training Accuracy :", r2_score(y_train, y_pred_train_ridge))
         print()
         print("Test Accuracy :", r2_score(y_test, y_pred_test_ridge))
         Ridge Model : [[ 8.14135029e-01 -3.00295914e-02 2.51212276e-02 1.37464222e+02
            2.88599737e+02]]
         Training Accuracy : 0.9394005370541024
         Test Accuracy: 0.9916566167602463
 In [ ]:
```

```
In []:
In []:
In [63]: sns.regplot(x=y_test, y=y_pred_test_ridge)
Out[63]: <Axes: xlabel='Profit'>
```



Elastic Net

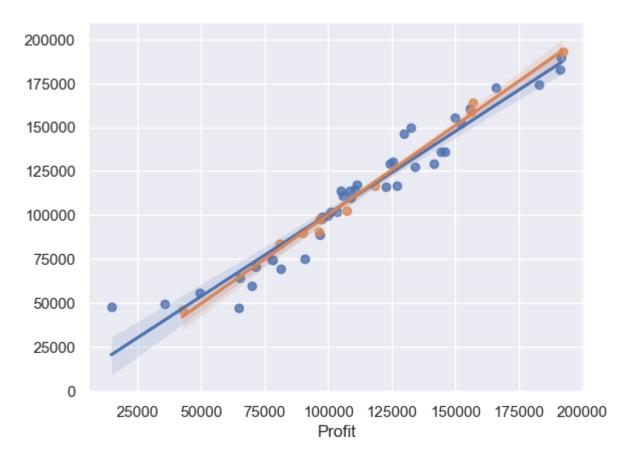
```
In [64]: elastic = ElasticNet(alpha=0.3, l1_ratio=0.1)
    elastic.fit(x_train, y_train)
    y_pred_train_elastic = elastic.predict(x_train)
    y_pred_test_elastic = elastic.predict(x_test)
    print("Training Accuracy :", r2_score(y_train, y_pred_train_elastic))
    print()
    print("Test Accuracy :", r2_score(y_test, y_pred_test_elastic))
```

Training Accuracy : 0.9393963701399531

Test Accuracy : 0.9916606769091509

```
In [65]: sns.regplot(x=y_train, y=y_pred_train_elastic)
sns.regplot(x=y_test, y=y_pred_test_elastic)
```

Out[65]: <Axes: xlabel='Profit'>



```
In [ ]: [
```

Gradient Descent

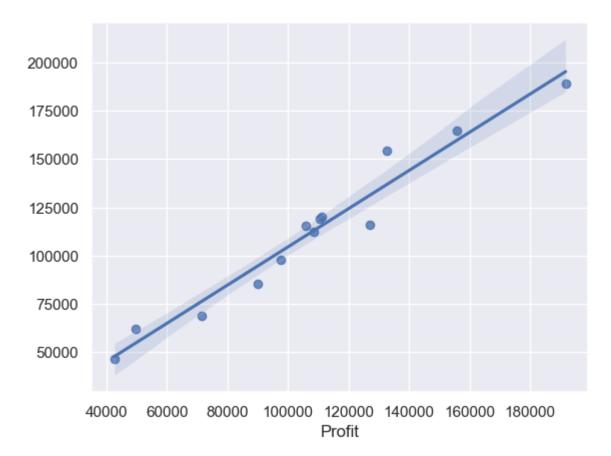
· Always scaled data is used.

```
In [66]: from sklearn.model_selection import train_test_split
         x_train, x_test, y_train, y_test = train_test_split(sc_x, y, test_size=0.25, random_state=101)
         print('x_train shape :',x_train.shape,'\n''x_test shape :', x_test.shape,'\n' 'y_train shape :',y_train.shape
         print()
         gd = SGDRegressor()
         gd.fit(x_train, y_train)
         y_pred_gd_train = gd.predict(x_train)
         y_pred_gd_test = gd.predict(x_test)
         Train_accuracy_gd =r2_score(y_train, y_pred_gd_train)
         Test_accuracy_gd = r2_score(y_test, y_pred_gd_test)
         print("GD Trainging Accuracy :", Train_accuracy_gd )
                                      :",Test_accuracy_gd )
         print("GD Test Accuracy
         x_train shape : (37, 5)
         x_test shape : (13, 5)
         y_train shape : (37, 1)
         y_test shape : (13, 1)
         GD Trainging Accuracy : 0.9473415978433063
         GD Test Accuracy
                             : 0.9418237075644532
```

The Gradient Descent is a best fit

```
In [67]: sns.regplot(x=y_test, y=y_pred_gd_test)
```

Out[67]: <Axes: xlabel='Profit'>



Ploynomial Regression

```
In [68]: poly = PolynomialFeatures()
    x_train_trans = poly.fit_transform(x_train)
    x_test_trans = poly.fit_transform(x_test)
    lr = LinearRegression()
    lr.fit(x_train_trans, y_train)
    poly.fit(x_train, y_train)
    y_pred_poly_train = lr.predict(x_train_trans)
    y_pred_poly_test = lr.predict(x_test_trans)
    Train_accuracy_poly = r2_score(y_train, y_pred_poly_train)
    Test_accuracy_poly = r2_score(y_test, y_pred_poly_test)

print("Polynomial Regression Trainging Accuracy :", Train_accuracy_poly )
    print("Polynomial Regression Test Accuracy :", Test_accuracy_poly )
```

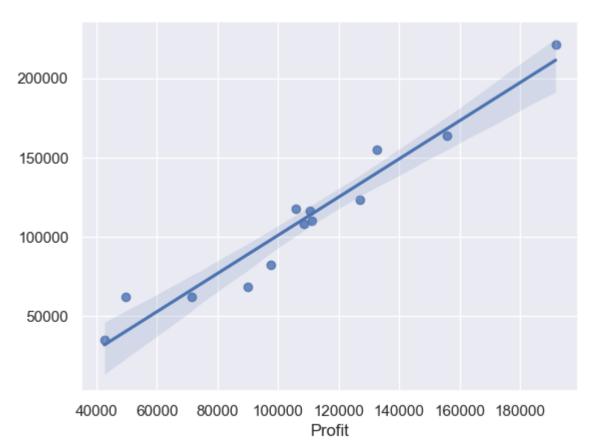
Polynomial Regression Trainging Accuracy : 0.9643335626660127 Polynomial Regression Test Accuracy : 0.8670425426922718

```
In [ ]:

In [ ]:
```

In [69]: sns.regplot(x=y_test, y=y_pred_poly_test)

Out[69]: <Axes: xlabel='Profit'>



Summary of the Model Building

- Test accuracy_MLR (test sample 20%) > Test accuracy_MLR (test sample 30%)
- Train accuracy MLR (test sample 20%) < Train accuracy MLR (test sample 30%)
- · Reguralization method enhanced the accuracy of both Train and Test

Out[93]:

	Test_size	R2_Train_Accuracy	R2_Test_Accuracy
MLR	0.3	0.965	0.899
OLS	0.3	0.993	0.983
Lasso	0.3	0.965	0.899
Ridge	0.3	0.960	0.900
ElasticNet	0.3	0.963	0.917
Gradient Descent	0.3	0.947	0.940
PolynomialRegression	0.3	0.964	0.867
MLR	0.2	0.932	0.983
OLS	0.2	0.987	0.995
Lasso	0.2	0.939	0.991
Ridge	0.2	0.939	0.991
ElasticNet	0.2	0.939	0.991
GradientDescent	0.2	0.947	0.941
Polynomial regression	0.2	0.964	0.867

```
In [98]: accuracy_df['variance']= abs(accuracy_df['R2_Train_Accuracy'] - accuracy_df['R2_Test_Accuracy'])*100
accuracy_df
```

Out[98]:

	Test_size	R2_Train_Accuracy	R2_Test_Accuracy	variance
MLR	0.3	0.965	0.899	6.6
OLS	0.3	0.993	0.983	1.0
Lasso	0.3	0.965	0.899	6.6
Ridge	0.3	0.960	0.900	6.0
ElasticNet	0.3	0.963	0.917	4.6
Gradient Descent	0.3	0.947	0.940	0.7
PolynomialRegression	0.3	0.964	0.867	9.7
MLR	0.2	0.932	0.983	5.1
OLS	0.2	0.987	0.995	0.8
Lasso	0.2	0.939	0.991	5.2
Ridge	0.2	0.939	0.991	5.2
ElasticNet	0.2	0.939	0.991	5.2
GradientDescent	0.2	0.947	0.941	0.6
Polynomial regression	0.2	0.964	0.867	9.7

Observation of the Model Building based on test sample size (0.2, 0.3)

- 1. Polynomial Regression is not a suitable model as the variance is quite high, it overfits the data. Which infers there is no non linearity between target and predictors.
- 2. The variance is minimum in case of Gradient descent.
- 3. The accuracy is getting better by reducing sample size.
- 4. In case of smaller sample size, the reguralization is not essential. MLR model is a best fit model for test sample size= 0.2



In []: